

Group Assignment 1

Joshua Jenkins, Peter Nguyen, Grace Hsieh,
Pablo Martinez Castro

Algorithm Design: Problem 1

- Matrix Multiplication for Square (nxn) Matrices
 - 1D Layout


$A(n \times n)$

\times

$B(n \times n)$

$=$

$C(n \times n)$



Results: Runtime



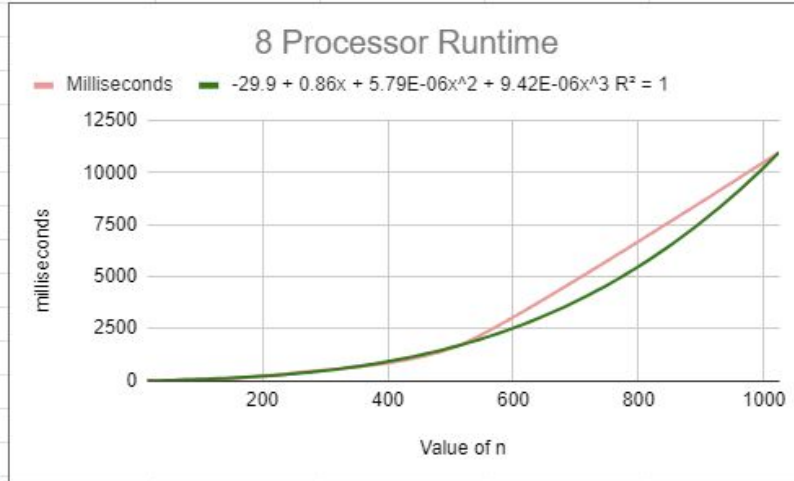
Run time (seconds)	P = 1	P = 2	P = 4	P = 8
n = 100	0.003	3.123	0.268	0.0656*
n = 1000	3.572	2.602	1.384	1.256
n = 5000	778.339	305.047	560.1	1181.91**
n = 10000	5618.92	2415.694	5958.35	9429.149**

*Calculated Value; $100 \% 8 \neq 0$

**HPC Context Switching

8 Processor Runtime Analysis - Estimation

n	milliseconds
8	0.618517
16	1.83251
32	3.76017
64	16.3313
128	58.191
256	387.012
512	1667.43
1024	10976.3



8 Processors			
nxn	Milliseconds	Seconds	Minutes
100	65.5779	<1 Second	<1 Minute
1000	10255.89	10.25589	<1 Minute
5000	1181914.85	1181.91485	19.69858083
10000	9429149.1	9429.1491	157.152485

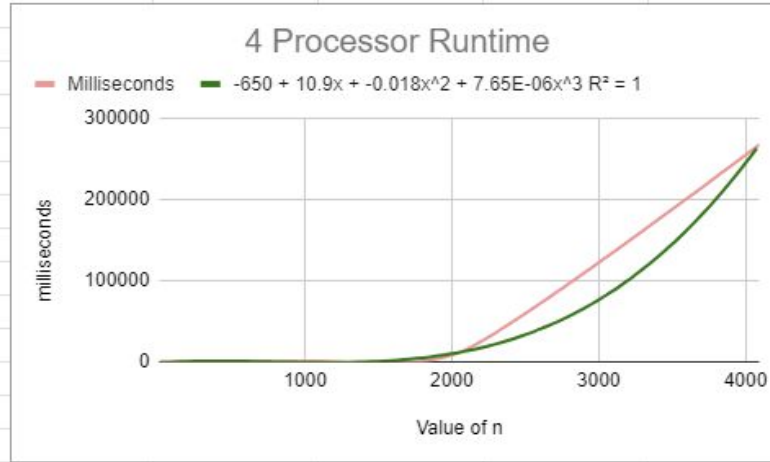
Based on:

$$-29.9 + 0.86x + 5.79E-06x^2 + 9.42E-06x^3$$

(line of best fit)

4 Processor Runtime Analysis - Estimation

n	milliseconds
4	0.05146
8	0.027399
16	0.032998
32	0.067823
64	0.363247
128	4.49998
256	23.6028
512	214.537
1024	1378.9
2048	11427.9
4096	267636

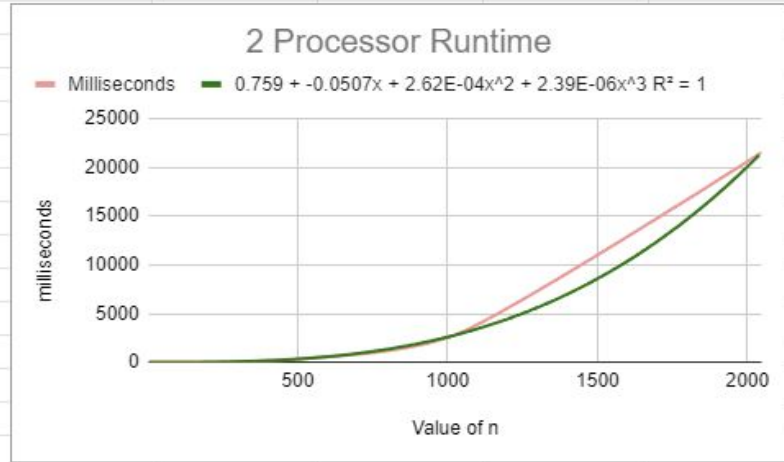


4 Processors			
nxn	Milliseconds	Seconds	Minutes
100	267.65	0.26765	<1 Minute
1000	1324.320484	1.324320484	<1 Minute
5000	560100	560.1	9.335
10000	5958350	5958.35	99.30583333

Based on:
 $-650 + 10.9x - 0.018x^2 + 7.65E-06x^3$
 (line of best fit)

2 Processor Runtime Analysis - Estimation

n	milliseconds
2	0.017772
4	0.009123
8	0.009576
16	0.016966
32	0.073625
64	0.536466
128	5.13548
256	40.0838
512	367.397
1024	2789.51
2048	21527
4096	530133



2 Processors


nxn	Milliseconds	Seconds	Minutes
100	3.123411375	<1 Second	<1 Minute
1000	2602.059	2.602059	0.04336765
5000	305047.259	305.047259	5.084120983
10000	2415693.759	2415.693759	40.26156265

Based on:

$0.759 + -0.0507x + 2.62E-04x^2 + 2.39E-06x^3$
(line of best fit)

Sequential Runtime Analysis - Estimation

n	milliseconds
100	3.18
1000	3572.38
5000	778339
10000	5572350



1 Processor					
nxn	Milliseconds	Seconds	Minutes	Hours	
100	3.18	<1 Second	<1 Minute	<1 Hour	
1000	3572.38	3.57	<1 Minute	<1 Hour	
5000	778339	778.34	12.97	<1 Hour	
10000	5572350	5572	92.87	1.55	

Based on:

$$20303 + -37.4x + 0.0163x^2 + 4.29E-06x^3$$

(line of best fit)

Results: Speedup

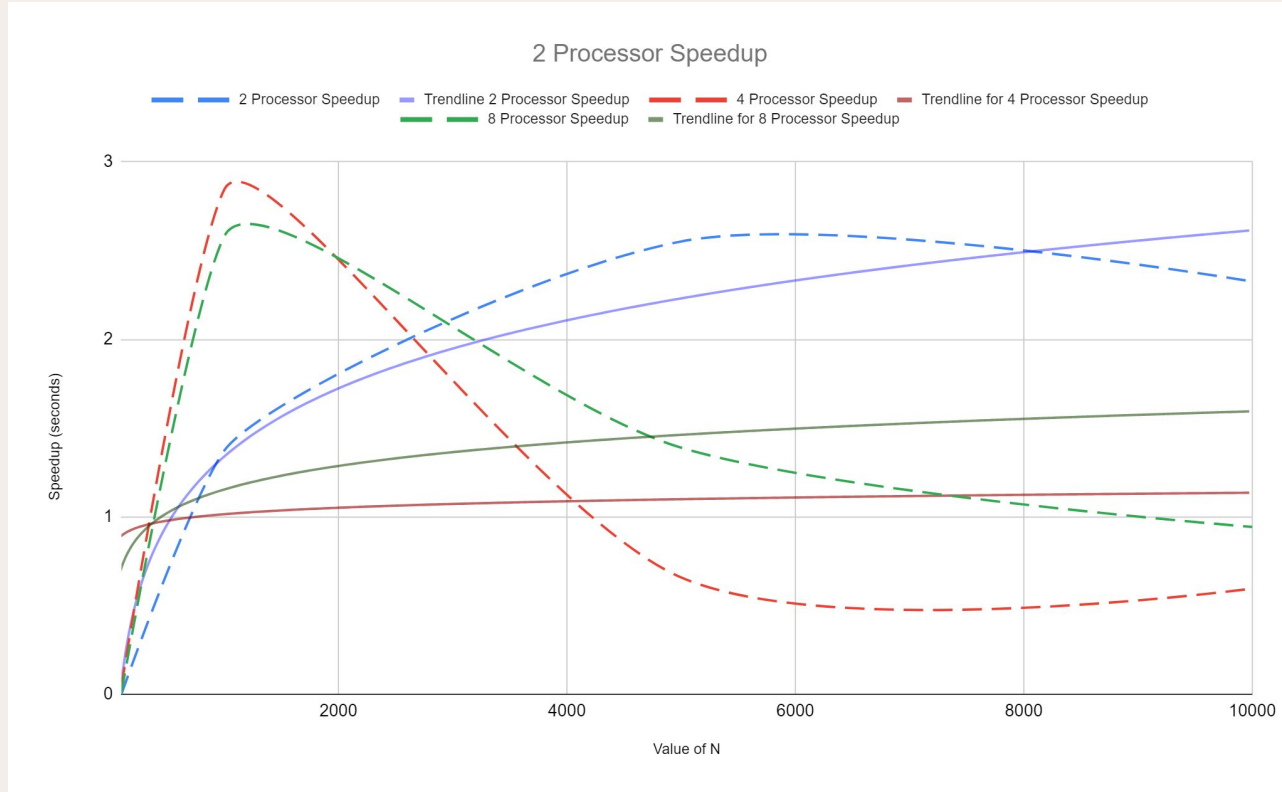


Speedup = Sequential Time / Parallel Time

Speedup (S)	P = 2	P = 4	P = 8
n = 100	0.00096	0.0112	0.0457
n = 1000	1.373	2.581	2.844
n = 5000	2.551	1.390	0.659*
n = 10000	2.326	0.943	0.596*

*HPC Context Switching

Results: Speedup (cont.)



Error Analysis

1. HPC context switching

- One error we saw occur was **context switching**.
- This caused our program to produce incorrect runtimes for some of our processes.
- This is due to multiple people using the HPC at the same time.

2. Line of best fit not accurate

- Another error we experienced is the calculation of runtimes not being calculated completely accurately after information was gathered.
- This was due to using the trendline equations to calculate runtime at certain values of n .

3. Potential Human Error

- Lastly, human error could have occurred. Recognizing and mitigating human error was essential for ensuring the validity of our computational analyses.

Algorithm Design: Problem 2

- $A(m \times n)$, $B(n \times q)$, $C(m \times q)$ - where $m \neq n \neq q$ (different sized matrices).
- **Data Distribution**
 - Divide A and B into non-uniform segments, based on their sizes and the number of processors
 - If m = rows of A, q = columns of B, and p = number of processors, each processor would handle m/p rows and q/p columns
 - Processor 0 will handle a larger portion of A and B, will also calculate the final result
 - Each processor p (excluding processor 0), will receive a portion of A and B
 - If the number of rows or column is not divisible evenly by the number of processors:
 - Distribute the remainder

Algorithm Design: Problem 2 Cont.



Algorithm Design:

1. Initialize the MPI environment.
 - `INITIALIZE MPI;`
2. Determine the total number of MPI processes and get the ID of the current process.
 - `DETERMINE MPI_WORLD_SIZE (p_total) AND RANK (pid);`
3. Root process initializes matrices $A[m][n]$ and $B'[n][q]$ (where B' is the transpose of $B[n][q]$).
 - `IF PID == 0 THEN INITIALIZE A(m,n), [B' ← TRANSPOSE(B[n][q])];`
4. Distribute matrix segments to all processes.
 - `localA, localB', p_total, pid ← DISTRIBUTE_MATRICES(A,B');`
5. Perform 1D matrix multiplication on the different segments.
 - `localC ← MM_1D_DISTRIBUTED(localA,localB',localM,localN,localQ);`
6. Gather all processes and assemble matrix $C[m][q]$
 - `C ← ASSEMBLE_MATRIX(GATHER_RESULTS(ALL localC's ACROSS ALL pid));`



Problem 2: DISTRIBUTE_MATRICIES ()

```
FUNCTION DISTRIBUTE_MATRICIES(in[A, B'] -> out[m, n, q, p_total, pid])
  IF pid == 0 THEN
    // Root process sends data to each process
    FOR process FROM 1 TO p_total - 1 DO
      SEND SEGMENT OF A TO process
      SEND SEGMENT OF B' TO process
    END FOR
  ELSE
    // Processes receive their information
    local_A <- RECEIVE SEGMENT OF A
    local_B' <- RECEIVE SEGMENT OF B'
  END IF
END FUNCTION
```

Communication Cost:

$$O(p_{total} * t_s + (m * n) * t_w)$$



Problem 2: MM_1D_DISTRIBUTED()

```
FUNCTION MM_1D_DISTRIBUTED(in[localA, localB'] -> out[localM, localN, localQ])
  localC <- INITIALIZE_MATRIX(localM, localQ)
  FOR i FROM 0 TO localM - 1 DO
    FOR j FROM 0 TO localQ - 1 DO
      sum <- 0
      FOR k FROM 0 TO localN - 1 DO
        sum <- sum + (localA[i][k] * localB'[j][k])
      END FOR
      localC[i][j] <- sum
    END FOR
  END FOR
  RETURN localC
END FUNCTION
```



Problem 2: GATHER_RESULTS()

```
FUNCTION GATHER_RESULTS(in[localC, m, q] -> out[C, p_total, pid])
  IF pid == 0 THEN
    // Root process gathers all segments of C
    C <- INITIALIZE_MATRIX(m, q)
    FOR process FROM 0 TO p_total - 1 DO
      RECEIVE SEGMENT OF C FROM process
    END FOR
  ELSE
    // Processes send result back to root processor
    SEND localC TO pid 0
  END IF
END FUNCTION
```

Communication Cost:

$$O(p_{total} * t_s + (m * q) * t_w)$$

Communication Cost & Runtime Analysis

$$\begin{aligned}\text{Overall Communication Costs} &= O(p_{total} * t_s + (m * n) * t_w) + O(p_{total} * t_s + (m * q) * t_w) \\ &= O(p_{total} * t_s + (m * (n + q)) * t_w)\end{aligned}$$

$$\text{Overall Runtime} = O\left(\frac{m * n * q}{p_{total}} + p_{total} * t_s + (m * (n + q)) * t_w\right)$$



Thank you!

